

Stats 202 Homework 9

Carina Zhang

12/7/2018


Citation: Code from labs in textbook.

Problem 1

Chapter 9, Exercise 2 (p. 368).


1 (a)

```
r = 2
plot(NA, NA, type = "n", xlim = c(-5, 5), ylim = c(-5, 5), asp = 1, xlab = "X1", ylab = "X2")
symbols(c(-1), c(2), circles = c(r), add = TRUE, inches = FALSE)
```




1 (b)

```
r = 2
plot(NA, NA, type = "n", xlim = c(-5, 5), ylim = c(-5, 5), asp = 1, xlab = "X1", ylab = "X2")
symbols(c(-1), c(2), circles = c(r), add = TRUE, inches = FALSE)
text(c(-1), c(2), "< 4")
text(c(-4), c(2), "> 4")
```



1 (c)

```
r = 2
plot(c(0, -1, 2, 3), c(0, 1, 2, 8), col = c("blue", "red", "blue", "blue"), type = "p", asp = 1, xlab = "X1", ylab = "X2",
      symbols(c(-1), c(2), circles = c(r), add = TRUE, inches = FALSE))
```



(0, 0), (2, 2), (3, 8) are classified to the blue class. (-1, 1) is classified to the red class.

1 (d)


We can expand the equation of the boundary, Since $(1+X_1)^2 + (2-X_2)^2 = 4$, we get $X_1^2 + X_2^2 + 2X_1 - 4X_2 + 1 = 0$. Thus, it is linear in terms of X_1, X_1^2, X_2, X_2^2 .

Problem 2

Chapter 9, Exercise 4 (p. 369).

```
set.seed(1)
x <- rnorm(100)
y <- 2*x^5 + x^6 + 4 + rnorm(100)
train <- sample(100, 50)
y[train] <- y[train] + 3
y[-train] <- y[-train] - 3

# Plot
plot(x[train], y[train], col = "green", xlab = "X", ylab = "Y")
points(x[-train], y[-train], col = "red")
```




SVM - Training

```
library(e1071)
set.seed(2)
z <- rep(0, 100)
z[train] <- 1

data <- data.frame(x = x, y = y, z = as.factor(z))
train <- sample(100, 50)
data.train <- data[train, ]
data.test <- data[-train, ]
```

```
svm.linear <- svm(z ~ ., data = data.train, kernel = "linear", cost = 10)
plot(svm.linear, data.train)
```

SVM classification plot



```
table = table(predict(svm.linear, data.train), data.train$z) #confusion matrix
table
```

```
##
##      0  1
##  0  0  0
##  1 20 30
```

```
error = (table[1,2] + table[2, 1]) / (table[1,1] + table[2, 2] + table[1,2] + table[2, 1])
error
```

```
## [1] 0.4
```


```
cat("The Training Error Rate is: ", error)
```

```
## The Training Error Rate is:  0.4
```

SVM - Poly - Training

```
set.seed(4)
svm.poly <- svm(z ~ ., data = data.train, kernel = "polynomial", cost = 10)
plot(svm.poly, data.train)
```

SVM classification plot



```
table = table(predict = predict(svm.poly, data.train), truth = data.train$z)
table

##      truth
## predict  0  1
##      0  1  0
##      1 19 30

error = (table[1,2] + table[2, 1]) / (table[1,1] + table[2, 2] + table[1,2] + table[2, 1])
error


## [1] 0.38
cat("The Training Error Rate is: ", error)

## The Training Error Rate is:  0.38
```

Third, we fit a support vector machine with a radial kernel.

```
svm.radial <- svm(z ~ ., data = data.train, kernel = "radial", cost = 20)
plot(svm.radial, data.train)
```

SVM classification plot



```
table = table(predict = predict(svm.radial, data.train), truth = data.train$z)
table

##      truth
## predict 0  1
##      0 20  1
##      1  0 29

error = (table[1,2] + table[2, 1]) / (table[1,1] + table[2, 2] + table[1,2] + table[2, 1])
error

## [1] 0.02
cat("The Training Error Rate is: ", error)
```


The Training Error Rate is: 0.02

We see that from all the training data error, both the polynomial-kernal (0.38) and Radial kernel(0.02) outperforms the SVM classifier (0.4).

SVM Linear - Test

```
plot(svm.linear, data.test)
```

SVM classification plot



```
table = table(predict = predict(svm.linear, data.test), truth = data.test$z)
table
```

```
##      truth
## predict 0  1
##      0  0  0
##      1 30 20
```

```
error = (table[1,2] + table[2, 1]) / (table[1,1] + table[2, 2] + table[1,2] + table[2, 1])
error
```

```
## [1] 0.6
```


```
cat("The Test Error Rate is: ", error)
```

```
## The Test Error Rate is:  0.6
```

SVM - poly - test

```
plot(svm.poly, data.test)
```

SVM classification plot



```
table = table(predict = predict(svm.poly, data.test), truth = data.test$z)
table
```

```
##      truth
## predict 0  1
##      0  0  0
##      1 30 20
```

```
error = (table[1,2] + table[2, 1]) / (table[1,1] + table[2, 2] + table[1,2] + table[2, 1])
error
```

```
## [1] 0.6
```


```
cat("The Test Error Rate is: ", error)
```

```
## The Test Error Rate is:  0.6
```

SVM - radial -

```
plot(svm.radial, data.test)
```


SVM classification plot



```
table = table(predict = predict(svm.radial, data.test), truth = data.test$z)
table

##      truth
## predict 0  1
##      0 29  2
##      1  1 18

error = (table[1,2] + table[2, 1]) / (table[1,1] + table[2, 2] + table[1,2] + table[2, 1])
error

## [1] 0.06

cat("The Test Error Rate is: ", error)
```

The Test Error Rate is: 0.06

We see that the linear, polynomial and radial support vector machines has a 0.6, 0.6, and 0.06 test error, respectively. So, radial kernel is the best model in this setting, with a multiple non-linear terms in the new boundary.

Problem 3

Chapter 9, Exercise 5 (p. 369).

3 (a)


Generate a data set with $n = 500$ and $p = 2$, such that the observations belong to two classes with a quadratic decision boundary between them.

```
set.seed(2)
x1 <- runif(500) - 0.5
x2 <- runif(500) - 0.5
y <- 1 * (x1^2 - x2^2 > 0)
```

3 (b)

Plot the observations, colored according to their class labels. Your plot should display X_1 on the x-axis and X_2 on the y-axis.

```
plot(x1, x2, xlab = "X1", ylab = "X2", col = (3 - y))
```



3 (c)

Fit a logistic regression model to the data, using X_1 and X_2 as predictors.


```
fit <- glm(y ~ x1 + x2, family = "binomial")
summary(fit)
```

```
##
## Call:
## glm(formula = y ~ x1 + x2, family = "binomial")
##
## Deviance Residuals:
##    Min       1Q   Median       3Q      Max
## -1.271  -1.193   1.097   1.147   1.209
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.07138    0.08959   0.797   0.426
## x1          -0.03532    0.29825  -0.118   0.906
```

```
## x2          0.27548    0.30762    0.896    0.370
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 692.50  on 499  degrees of freedom
## Residual deviance: 691.67  on 497  degrees of freedom
## AIC: 697.67
##
## Number of Fisher Scoring iterations: 3
```

3 (d)

```
data <- data.frame(x1 = x1, x2 = x2, y = y)
probs = predict(fit, data, type = "response")
preds <- rep(0, 500)
preds[probs > 0.49] = 1
plot(data[preds == 1,]$x1, data[preds == 1,]$x2, col = "red", xlab = "X1", ylab = "X2")
points(data[preds == 0,]$x1, data[preds == 0,]$x2, col = "green")
```



The decision boundary between the green dots and the red dots is linear (the very flat line at the bottom of the graph)

3 (e)

Now fit a logistic regression model to the data using non-linear functions of X1 and X2 as predictors.

```
fit_n1 <- glm(y ~ poly(x1, 3) + poly(x2, 2) + I(x1 * x2), family = "binomial")
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```


```
summary(fit_nl)
```

```
##
## Call:
## glm(formula = y ~ poly(x1, 3) + poly(x2, 2) + I(x1 * x2), family = "binomial")
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.275e-03 -2.000e-08  2.000e-08  2.000e-08  9.567e-04
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      83.52    3032.70   0.028   0.978
## poly(x1, 3)1    -262.79    38079.44 -0.007   0.994
## poly(x1, 3)2   26835.91    837602.91  0.032   0.974
## poly(x1, 3)3   -255.24    26750.33 -0.010   0.992
## poly(x2, 2)1     594.00    38344.53  0.015   0.988
## poly(x2, 2)2  -27011.68    842660.05 -0.032   0.974
## I(x1 * x2)       87.45    13359.71  0.007   0.995
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 6.9250e+02  on 499  degrees of freedom
## Residual deviance: 3.5668e-06  on 493  degrees of freedom
## AIC: 14
##
## Number of Fisher Scoring iterations: 25
```

3 (f)

As below, the decision boundary is non-linear. The non-linear decision boundary is surprisingly very similar to the true decision boundary.


```
probs <- predict(fit_nl, data, type = "response")
preds <- rep(0, 500)
preds[probs > 0.49] <- 1
plot(data[preds == 1, ]$x1, data[preds == 1, ]$x2, col = "red", xlab = "X1", ylab = "X2")
points(data[preds == 0, ]$x1, data[preds == 0, ]$x2, col = "green")
```



3 (g)

As we can see from the graph, a linear kernel with low cost classifies all points to a single class.


```
data$y <- as.factor(data$y)
svm.fit = svm(as.factor(y) ~ x1 + x2, data, kernel = "linear", cost = 0.01)
preds= predict(svm.fit, data)
plot(data[preds == 1, ]$x1, data[preds == 1, ]$x2, col = "red", xlab = "X1", ylab = "X2")
points(data[preds == 0, ]$x1, data[preds == 0, ]$x2, col = "green")
```



3 (h)

Similar to the previous parts, the radial kernel estimates the decision boundary well and resembles the true decision boundary.

```
data$y <- as.factor(data$y)
svmnl.fit <- svm(y ~ x1 + x2, data, kernel = "radial", gamma = 1)
preds <- predict(svmnl.fit, data)
plot(data[preds == 1, ]$x1, data[preds == 1, ]$x2, col = "red", xlab = "X1", ylab = "X2")
points(data[preds == 0, ]$x1, data[preds == 0, ]$x2, col = "green")
```



3 (i) Comment

We can see that SVMs with non-linear kernel and logistic regression with interaction terms are equally very powerful in finding non-linear boundaries. Also, SVM with linear kernel and logistic regression without any interaction term does not do well in finding non-linear decision boundaries. However, there is some manual efforts involved in tuning or picking right interaction term when it is a non-linear model. When we have a large number of features, the feature selection might become arbitrary. However, we can also implement cross-validation technique to make sure that we choose the best parameter.

Problem 4

Chapter 9, Exercise 8 (p. 371).

4 (a) & (b)

Create a training set containing a random sample of 800 observations, and a test set containing the remaining observations. Fit a support vector classifier to the training data using cost=0.01, with Purchase as the response and the other variables as predictors. Use the summary() function to produce summary statistics, and describe the results obtained.

```
library(ISLR)
set.seed(1)
train <- sample(nrow(OJ), 800)
OJ.train <- OJ[train, ]
OJ.test <- OJ[-train, ]

svm.linear <- svm(Purchase ~ ., data = OJ.train, kernel = "linear", cost = 0.01)
summary(svm.linear)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = OJ.train, kernel = "linear",
##      cost = 0.01)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##      cost:   0.01
##   gamma:    0.05555556
##
## Number of Support Vectors: 432
##
## ( 215 217 )
##
##
## Number of Classes: 2
##
## Levels:
## CH MM
```

As described in the summary data, the SVM model uses 432 support vectors. 215 of the vectors are in the CH class and 217 of the vectors are in the MM class.

4 (c)

What are the training and test error rates?

```
train.pred <- predict(svm.linear, OJ.train)
table = table(OJ.train$Purchase, train.pred)
table

##      train.pred
##      CH  MM
## CH 439  55
## MM  78 228

error = (table[1,2] + table[2, 1]) / (table[1,1] + table[2, 2])
error

## [1] 0.1994003

cat("The Train Error Rate is: ", error, "\n")

## The Train Error Rate is: 0.1994003
```

```

test.pred <- predict(svm.linear, OJ.test)
table = table(OJ.test$Purchase, test.pred)
table

##      test.pred
##      CH  MM
## CH 141  18
## MM  31  80

error = (table[1,2] + table[2, 1]) / (table[1,1] + table[2, 2])
error

## [1] 0.2217195

cat("The Test Error Rate is: ", error)

## The Test Error Rate is:  0.2217195

```

4 (d)

Use the `tune()` function to select an optimal cost. Consider values in the range 0.01 to 10

```

set.seed(1)
tune.out <- tune(svm, Purchase ~ ., data = OJ.train, kernel = "linear", ranges = list(cost = 10^seq(-2,
summary(tune.out)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##      cost
## 0.03162278
##
## - best performance: 0.17375
##
## - Detailed performance results:
##      cost  error dispersion
## 1 0.01000000 0.17500 0.03996526
## 2 0.03162278 0.17375 0.03884174
## 3 0.10000000 0.17875 0.03821086
## 4 0.31622777 0.17625 0.03701070
## 5 1.00000000 0.17750 0.03717451
## 6 3.16227766 0.18000 0.03496029
## 7 10.00000000 0.18000 0.04005205

```

The best forming cost is `cost = 0.03162278`

4 (e)

Compute the training and test error rates using this new value for cost.

```

set.seed(1)

svm.linear <- svm(Purchase ~ ., data = OJ.train, kernel = "linear", cost = 0.03162278)

```



```

train.pred <- predict(svm.linear, OJ.train)
table = table(OJ.train$Purchase, train.pred)
table

##      train.pred
##      CH  MM
## CH 438  56
## MM  72 234

error = (table[1,2] + table[2, 1]) / (table[1,1] + table[2, 2])
error

## [1] 0.1904762

cat("The Train Error Rate is: ", error, "\n")

```

```

## The Train Error Rate is:  0.1904762
test.pred <- predict(svm.linear, OJ.test)
table = table(OJ.test$Purchase, test.pred)
table

```

```

##      test.pred
##      CH  MM
## CH 138  21
## MM  31  80

error = (table[1,2] + table[2, 1]) / (table[1,1] + table[2, 2])
error

## [1] 0.2385321

cat("The Test Error Rate is: ", error)

```

```

## The Test Error Rate is:  0.2385321

```

4 (f)

```

set.seed(1)

svm.rad <- svm(Purchase ~ ., data = OJ.train, kernel = "radial")

train.pred <- predict(svm.rad, OJ.train)
table = table(OJ.train$Purchase, train.pred)
table

##      train.pred
##      CH  MM
## CH 455  39
## MM  77 229

error = (table[1,2] + table[2, 1]) / (table[1,1] + table[2, 2] + table[1,2] + table[2, 1])
error

## [1] 0.145

cat("The Train Error Rate is: ", error, "\n")

## The Train Error Rate is:  0.145

```

```

test.pred <- predict(svm.rad, OJ.test)
table = table(OJ.test$Purchase, test.pred)
table

##      test.pred
##      CH  MM
## CH 141  18
## MM   28  83

error = (table[1,2] + table[2, 1]) / (table[1,1] + table[2, 2] + table[1,2] + table[2, 1])
error

## [1] 0.1703704

cat("The Test Error Rate is: ", error)

## The Test Error Rate is:  0.1703704

#tune the model
tune.out <- tune(svm, Purchase ~ ., data = OJ.train, kernel = "radial", ranges = list(cost = 10^seq(-2,
summary(tune.out)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##      cost
## 0.3162278
##
## - best performance: 0.17125
##
## - Detailed performance results:
##      cost    error dispersion
## 1 0.01000000 0.38250 0.05596378
## 2 0.03162278 0.37250 0.06341004
## 3 0.10000000 0.17875 0.04168749
## 4 0.31622777 0.17125 0.05001736
## 5 1.00000000 0.17500 0.04750731
## 6 3.16227766 0.18000 0.04830459
## 7 10.00000000 0.18250 0.04866267

svm.rad <- svm(Purchase ~ ., data = OJ.train, kernel = "radial", cost = 0.3162278)

train.pred <- predict(svm.rad, OJ.train)
table = table(OJ.train$Purchase, train.pred)
table

##      train.pred
##      CH  MM
## CH 448  46
## MM   78 228

error = (table[1,2] + table[2, 1]) / (table[1,1] + table[2, 2] + table[1,2] + table[2, 1])
error

## [1] 0.155

```

```

cat("The Train Error Rate is: ", error, "\n")

## The Train Error Rate is: 0.155
test.pred <- predict(svm.rad, OJ.test)
table = table(OJ.test$Purchase, test.pred)
table

##      test.pred
##      CH  MM
## CH 144  15
## MM  29  82

error = (table[1,2] + table[2, 1]) / (table[1,1] + table[2, 2] + table[1,2] + table[2, 1])
error

## [1] 0.162963
cat("The New Test Error Rate with best Cost is: ", error)

## The New Test Error Rate with best Cost is: 0.162963

```

4 (g)

```

set.seed(1)

svm.poly <- svm(Purchase ~ ., data = OJ.train, kernel = "polynomial", degree = 2)

train.pred <- predict(svm.poly, OJ.train)
table = table(OJ.train$Purchase, train.pred)
table

##      train.pred
##      CH  MM
## CH 461  33
## MM 105 201

error = (table[1,2] + table[2, 1]) / (table[1,1] + table[2, 2] + table[1,2] + table[2, 1])
error

## [1] 0.1725
cat("The Train Error Rate is: ", error, "\n")

## The Train Error Rate is: 0.1725
test.pred <- predict(svm.poly, OJ.test)
table = table(OJ.test$Purchase, test.pred)
table

##      test.pred
##      CH  MM
## CH 149  10
## MM  41  70

error = (table[1,2] + table[2, 1]) / (table[1,1] + table[2, 2] + table[1,2] + table[2, 1])
error

## [1] 0.188889

```

```

cat("The Test Error Rate is: ", error)

## The Test Error Rate is: 0.1888889
#tune the model
tune.out <- tune(svm, Purchase ~ ., data = OJ.train, kernel = "polynomial", ranges = list(cost = 10^seq(
summary(tune.out)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##      cost
## 3.162278
##
## - best performance: 0.175
##
## - Detailed performance results:
##      cost      error dispersion
## 1 0.01000000 0.36750 0.05041494
## 2 0.03162278 0.34625 0.05894029
## 3 0.10000000 0.30250 0.07402139
## 4 0.31622777 0.20250 0.06368324
## 5 1.00000000 0.18750 0.05368374
## 6 3.16227766 0.17500 0.05559027
## 7 10.00000000 0.18125 0.04686342

The best performing is cost = 3.162278

svm.poly <- svm(Purchase ~ ., data = OJ.train, kernel = "polynomial", cost = 3.162278)

train.pred <- predict(svm.poly, OJ.train)
table = table(OJ.train$Purchase, train.pred)
table

##      train.pred
##      CH  MM
## CH 460  34
## MM  76 230

error = (table[1,2] + table[2, 1]) / (table[1,1] + table[2, 2] + table[1,2] + table[2, 1])
error

## [1] 0.1375

cat("The New Train Error Rate with best cost is: ", error, "\n")

## The New Train Error Rate with best cost is: 0.1375

test.pred <- predict(svm.poly, OJ.test)
table = table(OJ.test$Purchase, test.pred)
table

##      test.pred
##      CH  MM
## CH 142  17
## MM  37  74

```

```
error = (table[1,2] + table[2, 1]) / (table[1,1] + table[2, 2] + table[1,2] + table[2, 1])
error
```

```
## [1] 0.2
```

```
cat("The New Test Error Rate with best Cost is: ", error)
```

```
## The New Test Error Rate with best Cost is:  0.2
```

4 (h)

In terms of training error, the tuned radial kernel SVM has the smallest training error.

In terms of testing error, the tuned polynomial kernel SVM has the smallest testing error

Overall, the non-linear SVM kernel seem to give the best results on this data with smaller training/test error.